



# An overview of the TurboFan compiler

Benedikt Meurer / [bmeurer@chromium.org](mailto:bmeurer@chromium.org) / [@bmeurer](https://twitter.com/bmeurer)

V8 Compiler Tech Lead  
Google Munich

# Overview

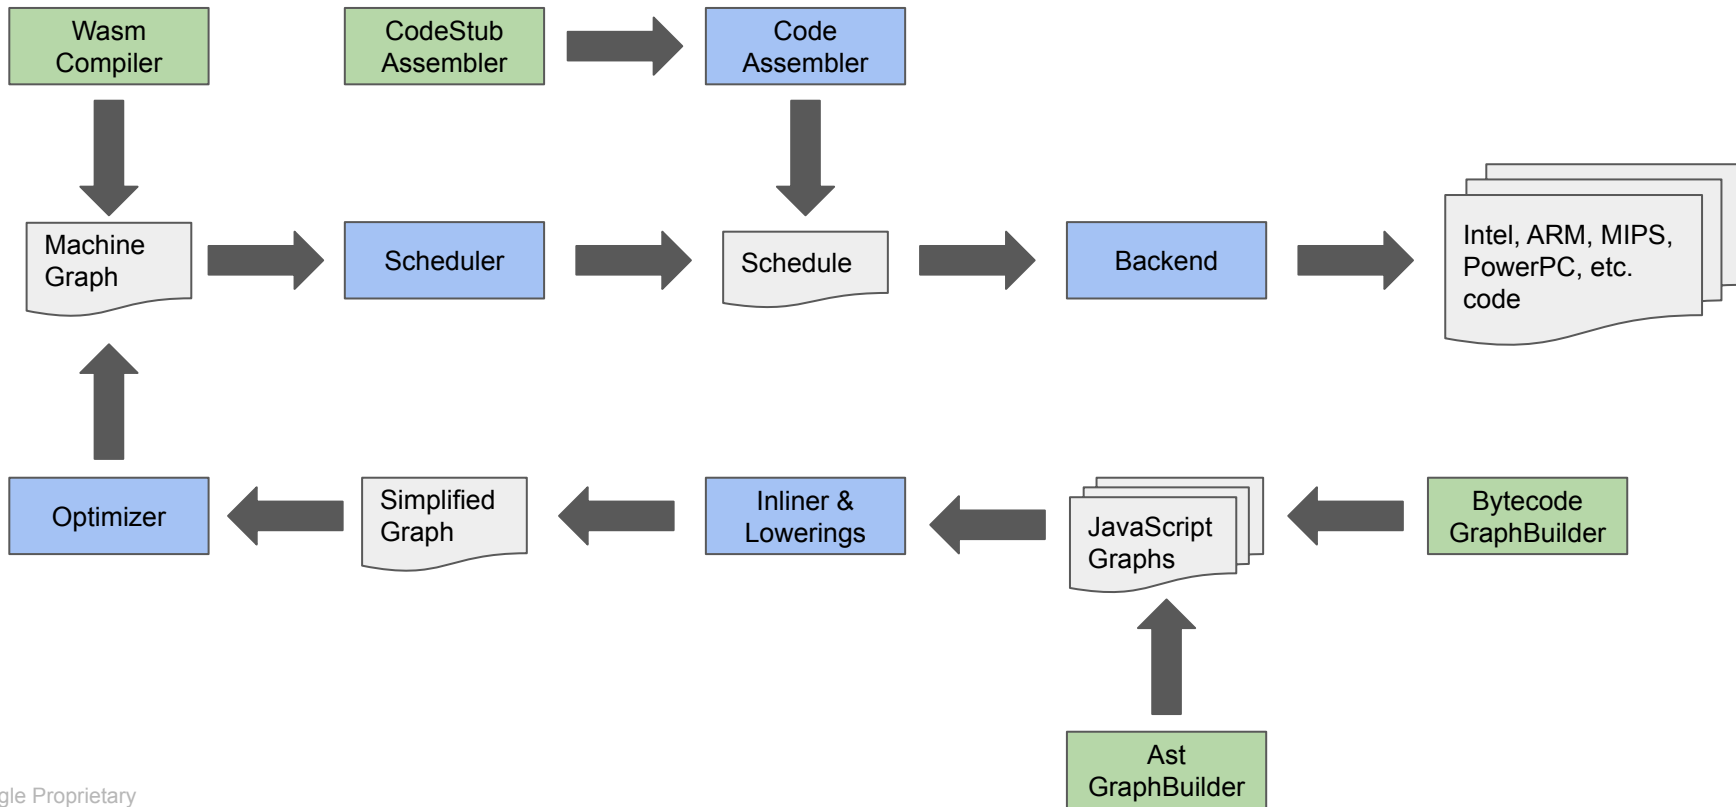
What is TurboFan?

- **the new** JavaScript compiler to replace the aging Crankshaft
- a high-level, portable assembler

Why a portable assembler?

- Interpreter bytecode handlers
- Code stubs and builtins
- asm.js & WebAssembly

# The TurboFan architecture / entry points



# TurboFan as a portable assembler

# Motivation

Traditionally V8 contains a lot of hand-written native code:

- Fast paths for EcmaScript builtins (i.e. `Object.prototype.toString`)
- IC handlers and dispatchers
- Code stubs (i.e. `FastNewClosureStub`)

Native code on the web efforts:

- `asm.js`
- WebAssembly

# Goals

- Provide a uniform code generation architecture
- Reduce porting / maintenance overhead of V8 (currently already 10 ports!)
- Remove performance cliffs due to slow builtins
- Make experimenting with new features easier (i.e. changes to load/store ICs, bootstrapping an interpreter)

# CodeAssembler vs. CodeStubAssembler

- CodeAssembler as shim to the TurboFan machine-level IR; wraps RawMachineAssembler, which generates a Schedule with machine instructions
- CodeStubAssembler adds high-level V8/JavaScript-related macros on-top
- InterpreterAssembler adds Ignition-related macros on-top of CodeStubAssembler
- CodeAssembler is portable assembler entry point for TurboFan

# Benefits

- Fewer performance cliffs
- Better baseline performance (real world benefit, i.e. `ForInFilter`, `Object.prototype.toString`, `Function.prototype.bind`, **etc.**)
- Fewer bugs (no more manual register allocation, etc.)
- Reduce porting overhead
- Refactoring viable



# TurboFan as JavaScript compiler

# Motivation

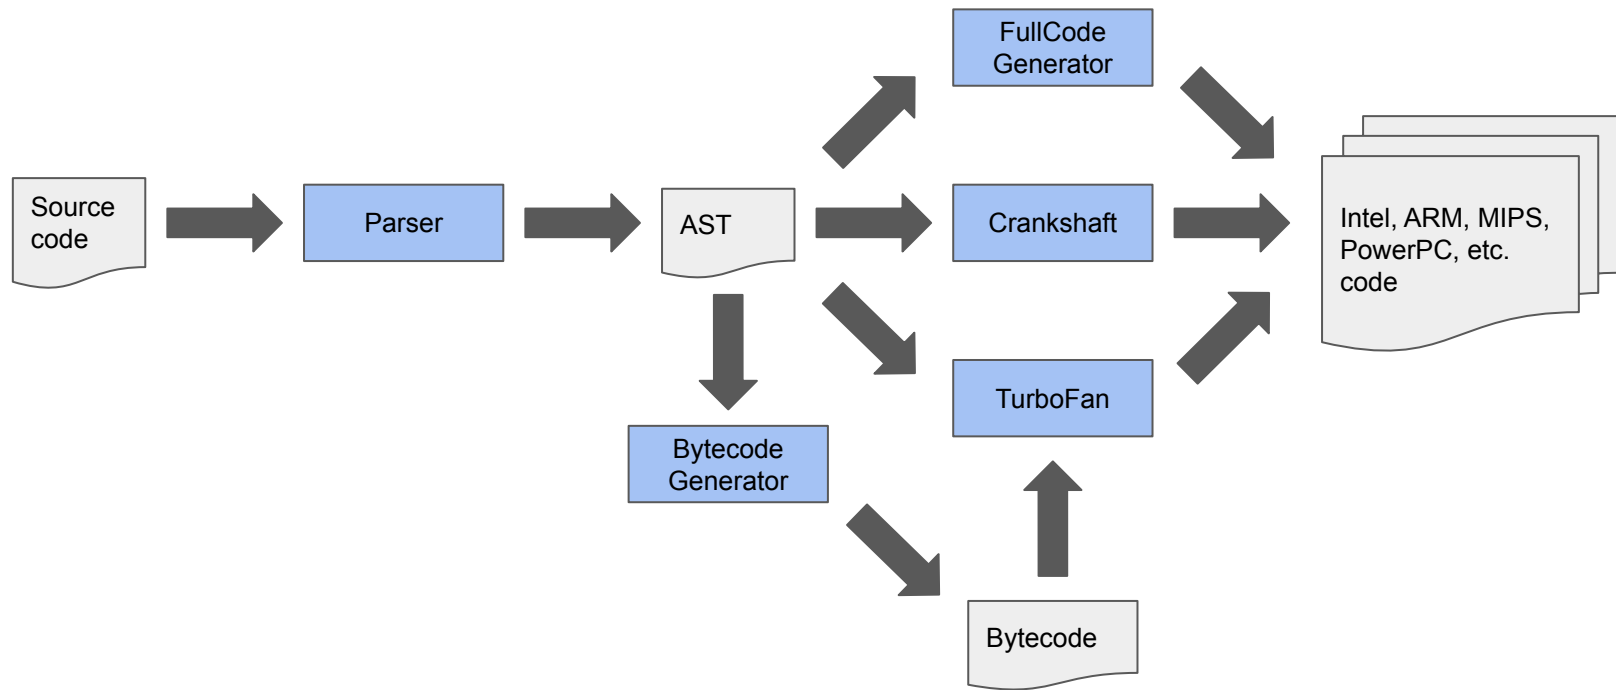
Crankshaft served us well, but has various shortcomings:

- Doesn't scale to full, modern JavaScript (`try-catch`, `for-of`, generators, `async/await`, ...)
- Defaults to deoptimization (performance cliffs, deoptimization loops)
- Graph construction, inlining and optimization all mixed up
- Tight coupling to fullcodegen / brittle environment tracking
- Limited optimization potential / limited static analysis (i.e. type propagation)
- High porting overhead
- Mixed low-level and high-level semantics of instructions

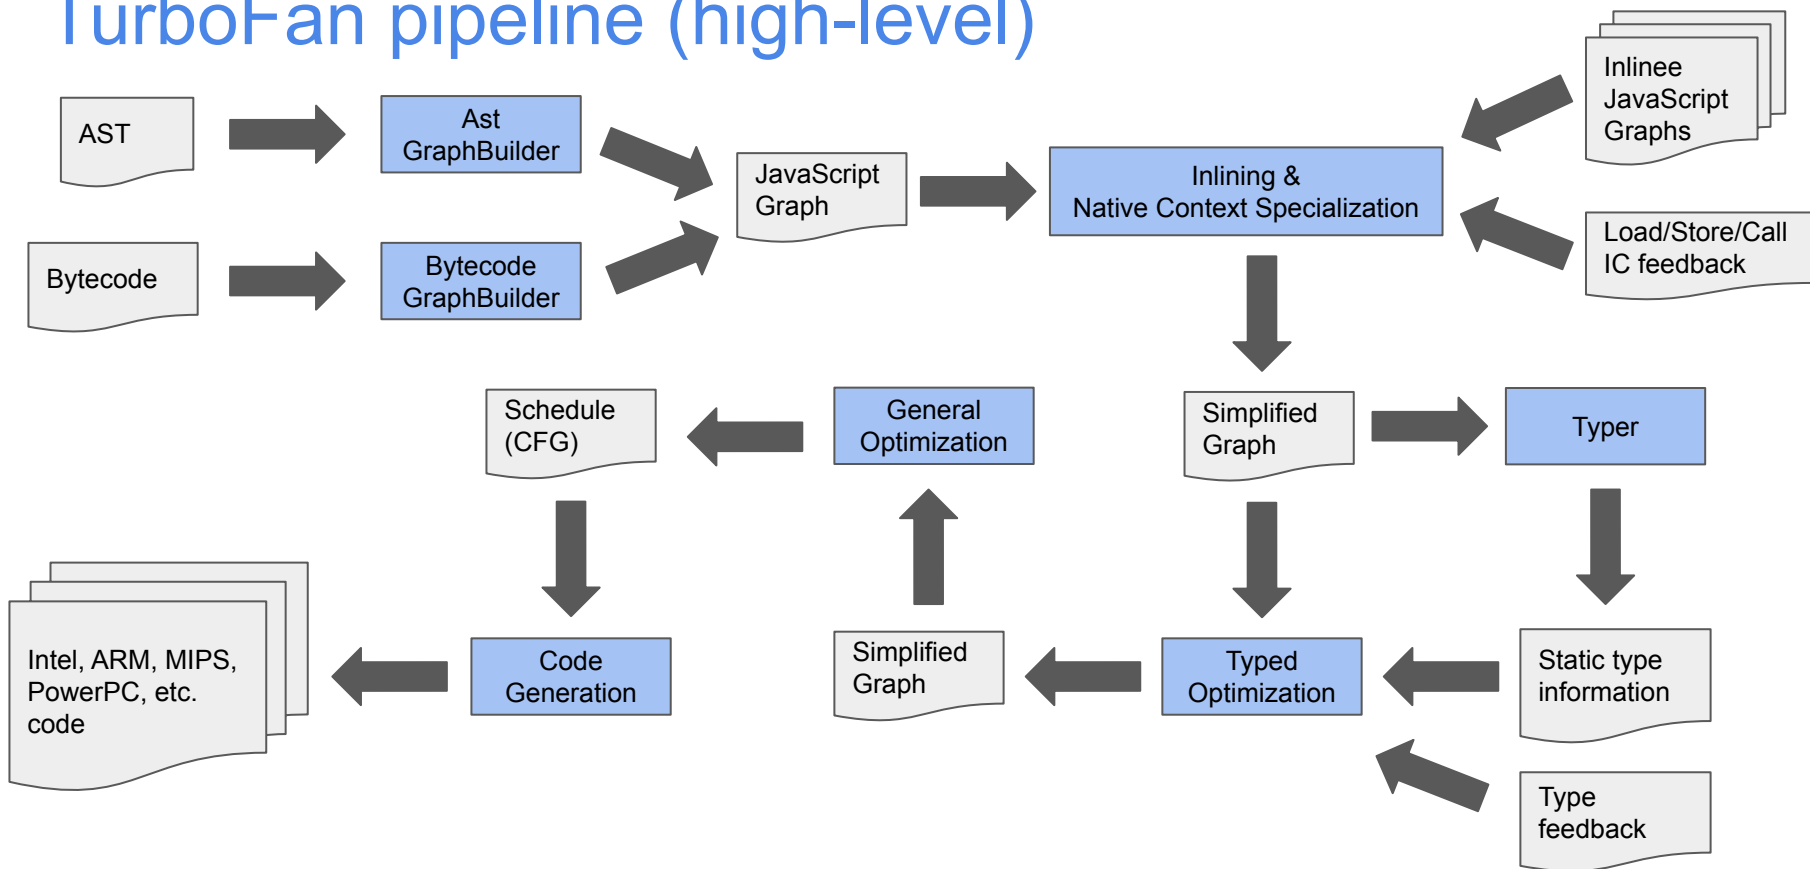
# Goals

- Full ESnext language support (`try-catch/-finally`, `class` literals, `eval`, `generators`, `async` functions, `modules`, `destructuring`, etc.)
- Utilize and propagate (static) type information
- Separate graph building from optimization / inlining
- No deoptimization loops / deoptimization only when really beneficial
- Sane environment tracking (also for lazy deoptimization)
- Predictable peak performance

# V8 compilation overview



# TurboFan pipeline (high-level)



# Graph Building

- Turns Bytecode or AST into graph with JavaScript nodes
  - JSAdd, JSCallFunction, JSLoadProperty, etc.
  - Branch, IfTrue, IfFalse, etc.
- Tracks environment for baseline code and inserts appropriate FrameStates
  - CheckPoint with FrameState for eager deoptimization
  - FrameState after nodes that can lead to arbitrary JS execution (i.e. everything except JSTypeOf, JSStrictEqual and JSToBoolean) for lazy deoptimization
- Eliminates dead FrameState uses, i.e. values that baseline cannot see are replaced with `optimized_out` sentinel
- Optional OSR deconstruction right afterwards to remove the non-OSR entry

# Native Context Specialization & Inlining

- Specialize global object property access to the global object for native context
- Specialize to native context based on Load/Store/Call IC feedback
  - JSLoadNamed[x](o) turns into fast access for x on o if baseline collected feedback
  - Specialize function calls to direct calls (i.e. to known JSFunctions) when Call IC provides feedback
- Inline through JSCallFunction and JSCallConstruct sites, even [polymorphic sites](#)
  - Supports [sane heuristics](#) (not only first come, first serve)
  - Can even inline through `Function.prototype.call/.apply`
- Eliminates dead nodes in the (control) graph on-the-fly
- Doesn't take into account type feedback for binary/compare operations

# Typer & Typed Optimization

- Typer assigns types to value producing nodes, based on JS typing rules
  - Already considers types for fields known from native context specialization
  - Propagates type information through phis and sigmas (loop exits)
  - Range analysis and constant folding (some) included
- Lowering JS graph to simplified graph based on types
  - Take into account *static* type information first, but also utilize type feedback (binary/compare)
  - Inline builtin calls based on parameter types (i.e. Math.abs, Array.prototype.push, etc.)
- Lower JS object creation nodes to inline allocations
  - JSCreate becomes an inline allocation when new.target is known
  - Similar for JSCreateArray, JSCreateLiteralObject, etc.
  - Makes those eligible for escape analysis and load/store/check elimination



# General Optimizations

- Loop peeling
  - No more deoptimization loops because of aggressive hoisting
- Load/check elimination
  - Kills redundant loads and checks
  - Includes global value numbering pass
- Escape analysis
  - Eliminate non escaping allocations
  - Scalar replacement of aggregates
  - Not yet partial; allocation sinking planned
- Representation selection
  - Truncation analysis
  - Type feedback propagation

# General Optimizations

- Effect/control linearization
  - Wires potentially deoptimizing nodes into effect/control chain, assigns check points
  - Expands macro operations with internal control flow (i.e. ChangeTaggedToFloat64)
  - Optimizes branches that depend onphis (branch cloning)
- Redundant store elimination
- Control flow optimization (turns certain branch chains into switches)
- Allocation folding and write barrier elimination
  - Fragmentation-free folding of allocations
  - Even across branches; not yet across merges
- Late optimization pass
  - Another global value numbering pass
  - Strength reduction, dead code elimination, redundant branch elimination

# Code Generation

- **Sophisticated instruction selection**
  - Pattern matching on subgraphs for often optimal cover
  - Architecture-dependent matchings
  - Dead instruction elimination
- **Register allocation**
  - Essentially improved version of the lithium allocator
- **Jump threading**
  - Remove jumps between blocks
- **Code emission**

# Code generation example

Generate code for `add`

```
function add(x, y) { return x + y; }
```

Run with `--turbo --ignition-staging`

```
add(1, 2);  
add(2, 3);  
%OptimizeFunctionOnNextCall(add);  
add(3, 4);
```

Almost optimal code sequence already (Smi representation missing)

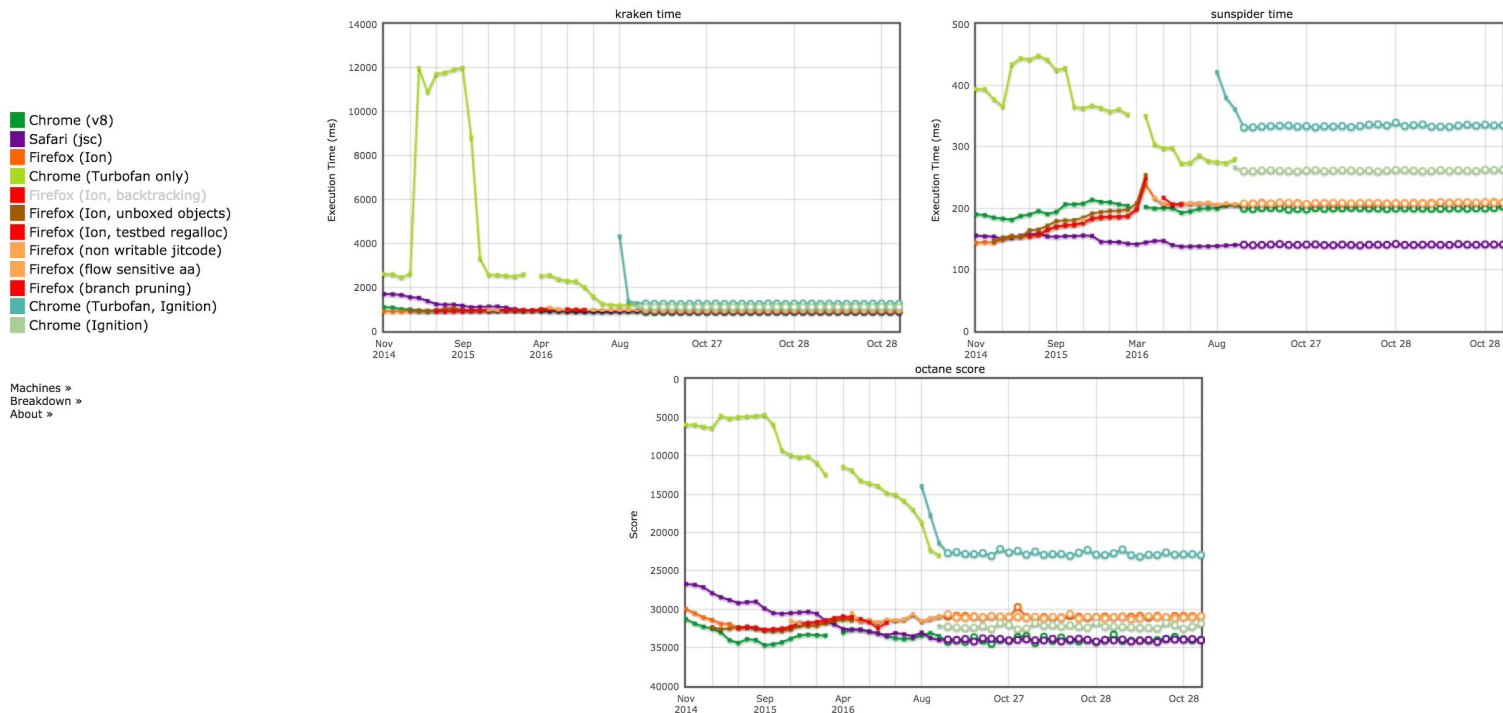
# Code generation example

```
-- B0 start (construct frame) --
0x3caafc104060  0 55          push rbp
0x3caafc104061  1 4889e5      REX.W movq rbp, rsp
0x3caafc104064  4 56          push rsi
0x3caafc104065  5 57          push rdi
0x3caafc104066  6 493ba5600c0000 REX.W cmpq rsp, [r13+0xc60]
0x3caafc10406d 13 0f863d000000  jna 80 (0x3caafc1040b0)
-- B2 start --
-- B3 start (deconstruct frame) --
0x3caafc104073 19 488b4518     REX.W movq rax, [rbp+0x18]
0x3caafc104077 23 a801        test al, 0x1
0x3caafc104079 25 0f8548000000 jnz 103 (0x3caafc1040c7)
0x3caafc10407f 31 488b5d10     REX.W movq rbx, [rbp+0x10]
0x3caafc104083 35 f6c301      testb rbx, 0x1
0x3caafc104086 38 0f8540000000 jnz 108 (0x3caafc1040cc)
0x3caafc10408c 44 488bd3      REX.W movq rdx, rbx
0x3caafc10408f 47 48c1ea20     REX.W shrq rdx, 32
0x3caafc104093 51 488bc8      REX.W movq rcx, rax
0x3caafc104096 54 48c1e920     REX.W shrq rcx, 32
0x3caafc10409a 58 03d1        addl rdx, rcx
0x3caafc10409c 60 0f802f000000 jo 113 (0x3caafc1040d1)
0x3caafc1040a2 66 48c1e220     REX.W shlq rdx, 32
0x3caafc1040a6 70 488bc2      REX.W movq rax, rdx
0x3caafc1040a9 73 488be5      REX.W movq rsp, rbp
0x3caafc1040ac 76 5d          pop rbp
0x3caafc1040ad 77 c21800      ret 0x18
-- B4 start (no frame) --
-- B1 start (deferred) --

0x3caafc1040b0 80 48bbd012abb3397f0000 REX.W movq rbx, 0x7f39b3ab12d0
0x3caafc1040ba 90 33c0        xorl rax, rax
0x3caafc1040bc 92 488b75f8     REX.W movq rsi, [rbp-0x8]
-- real.js:229:12 --
0x3caafc1040c0 96 e87b02f0ff    call 0x3caafc004340    ;; code: STUB
0x3caafc1040c5 101 ebac         jmp 19 (0x3caafc104073)
0x3caafc1040c7 103 e834ffd7ff    call 0x3caafbe84000    ;; deoptimization
bailout 0
0x3caafc1040cc 108 e839ffd7ff    call 0x3caafbe8400a    ;; deoptimization
bailout 1
0x3caafc1040d1 113 e83effd7ff    call 0x3caafbe84014    ;; deoptimization
bailout 2
0x3caafc1040d6 118 90        nop
0x3caafc1040d7 119 90        nop
0x3caafc1040d8 120 90        nop
0x3caafc1040d9 121 90        nop
0x3caafc1040da 122 90        nop
0x3caafc1040db 123 90        nop
0x3caafc1040dc 124 90        nop
0x3caafc1040dd 125 90        nop
0x3caafc1040de 126 90        nop
0x3caafc1040df 127 90        nop
0x3caafc1040e0 128 90        nop
0x3caafc1040e1 129 90        nop
0x3caafc1040e2 130 90        nop
0x3caafc1040e3 131 90        nop
;;; Safepoint table.
```

# Where do we stand?

# Peak-performance benchmarks



# What's coming?



# Future

- Retire AstGraphBuilder (asm.js only soon)
  - WebAssembly based asm validator coming
- Optimize ESnext features via Ignition+TurboFan
- Add missing optimizations to TurboFan (i.e. arguments object, Smi representation, etc.)
- Eventually route more (traditional) JavaScript through Ignition+TurboFan only

# Upcoming talks

- [TurboFan IR overview](mailto:jarin@chromium.org) ([jarin@chromium.org](mailto:jarin@chromium.org))
  - Different IR layers (JavaScript, simplified, machine)
  - Sea-of-nodes overview
- TurboFan deoptimization support ([jarin@chromium.org](mailto:jarin@chromium.org))
  - Check point chain
  - Lazy bailout
  - Deoptimizer support
  - Liveness analysis

# Questions?