



# TurboFan JIT Design

Ben L. Titzer  
Google Munich

## V8 Background

- JavaScript has some difficult to optimize features
  - No explicit types
  - Prototype-based property lookup
  - Dynamic evaluation of code
- V8 was the first really \*fast\* JavaScript VM
  - Sophisticated and efficient object layout
  - Compile-only: no interpreter
    - Quick, non-optimizing JIT (fullcode)
    - Inline caching, type feedback
    - Generational GC with low pause times
- V8 launched with Chrome in 2008
  - Optimizing JIT (CrankShaft) launched in 2010


# TurboFan Design Goals

- Achieve best peak performance
  - Highest quality machine code
  - Within normal constraints of JIT compilation
- Make best use of static type information
  - asm.js, latent JavaScript types, TypeScript, SoundScript proposal
- Reduce platform-specific implementation effort
  - Better separation between front, middle, and backend of compiler
- Improve testability
  - Prevent correctness bugs and verify optimizations activate

# TurboFan Program Representation (IR)


- NOT: Control Flow Graph (CFG)
  - Fully-specified evaluation order; e.g. pure operations like integer addition
- INSPIRATION: Sea of Nodes
  - Relax evaluation order for most operations
  - Effect edges order stateful operations
  - Skeleton of a CFG remains
  - Why? Better redundant code elimination, more code motion
- REALLY: “Soup” of Nodes
  - Relax Sea of Nodes control flow subgraph even further
  - Disconnected “floating control” islands offer more scheduling freedom
  - Why? Lowering of language levels, even more code motion

# Do not get seasick!




All computations are expressed as nodes in the sea of nodes


Edges represent dependencies between computations




## Dependencies constrain Ordering



 legal





 legal




 illegal


# Lack of Ordering means Compiler Freedom!



# Larger class of equivalences




graph  
creation





# The Sea is SSA


$x = 3 * 8$   
 $x + 3$



No such thing as local variables!

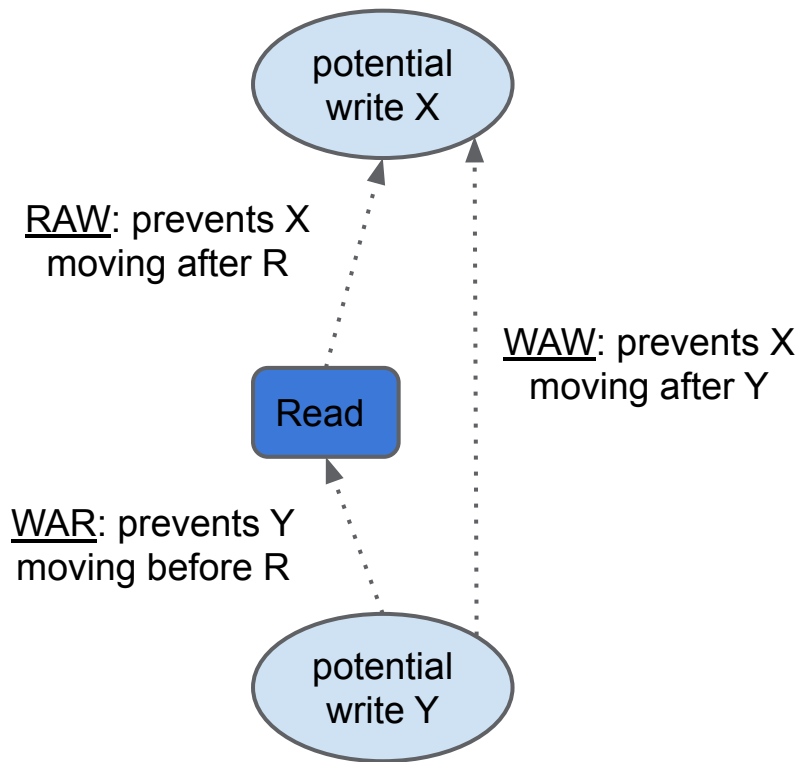
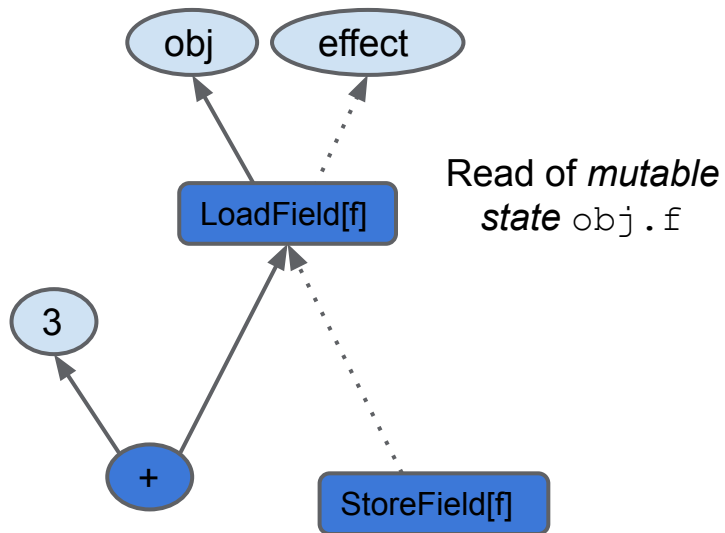
Graph building from source renames locals

SSA  
renaming



Multiple incoming  
edges possible

# Effect Edges




# Expressing Control

- Nodes: express computation
  - Constants, parameters, arithmetic, load, store, calls
  - Source program is SSA renamed so locals are substituted with nodes
- Edges: express dependencies (constrain order)
  - dataflow edges express using the value output of a computation
  - effect edges order operations reading and writing state
- NEXT: Control with start, branches, loops, merge, and end
  - How do we express non-straight line code?


# Control nodes and Control edges



straightline program




branch



while loop

# Our first complete graph

```
function (x) { return x + 3; }
```



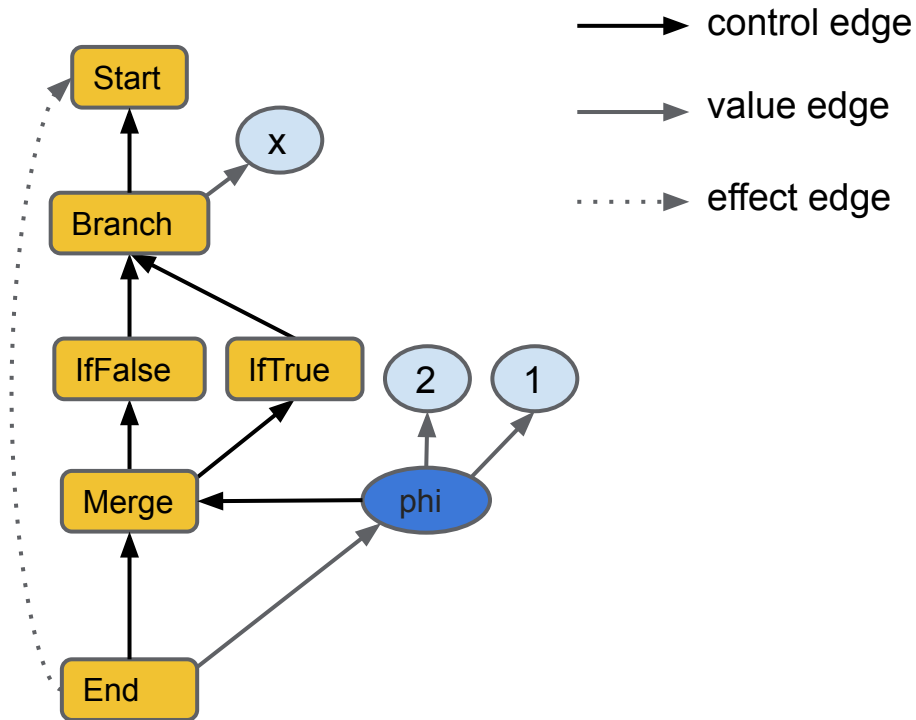
→ control edge

→ value edge

...→ effect edge

## Branch example

```
function (x) { return x ? 1 : 2; }
```




# Language Levels


- JavaScript: (“JS”) operators
  - Express semantics of JavaScript’s overloaded operators
  - Produce and consume effects in the graph
- Intermediate: (“Simplified”) operators
  - Express VM-level operations, such as allocation, bounds checks
  - Arithmetic independent of number representation
- Machine: (“Machine”) operators
  - Correspond closely to single machine instructions
  - Most have no side effects
  - Must be supported by backend for each platform

# Language Levels


- JavaScript:



- Intermediate:



- Machine:






# Type and Range Analysis

- JavaScript is not statically typed
  - Values have types, not variables
  - 8 is a Number, "x" is a String
  - All basic operators (+ - \* / % == !=) overloaded for objects
- All arithmetic is done in 64-bit floating point
  - Empirically, most programs use only small integers ( $\leq 31$ bits)
  - Overflow to double usually causes code to bailout to slow path
  - Troublesome cases: NaN, Infinity, -Infinity, -0.0
- asm.js language subset
  - Annotations such as  $(x + y) \mid 0$
  - Truncation maps NaN, Infinity, -Infinity, -0.0 to integer 0


# Type and Range Analysis



typing  
alone




**+: Num**



typing + range  
analysis




**+: Int in [11, 26]**



typing  
alone



**phi: Int**



typing + range  
analysis




**phi: Int in [0, 15]**


# Optimization

- Nearly all optimization happens on the sea of nodes
  - Top-down or bottom-up graph transformations
  - Isolates transformations from error-prone ordering of computations
  - Local reasoning leads to incremental transformations
- Reachability => Liveness
  - Nodes not reachable from end are *dead*
    - Including dead control, dead effects, dead computation
  - Most phases never see dead code
  - Dead code never placed in final schedule


# Reduction





constant  
folding




strength  
reduction




strength  
reduction




phi  
simplification





algebraic  
reassociation





# Typed Lowering as Reduction





typed  
lowering




typed  
lowering









typed  
lowering




generic  
lowering




# Global Value Numbering as Reduction






# Control Optimization as Reduction




# Control Optimization as Reduction



merge  
reduction




merge  
reduction

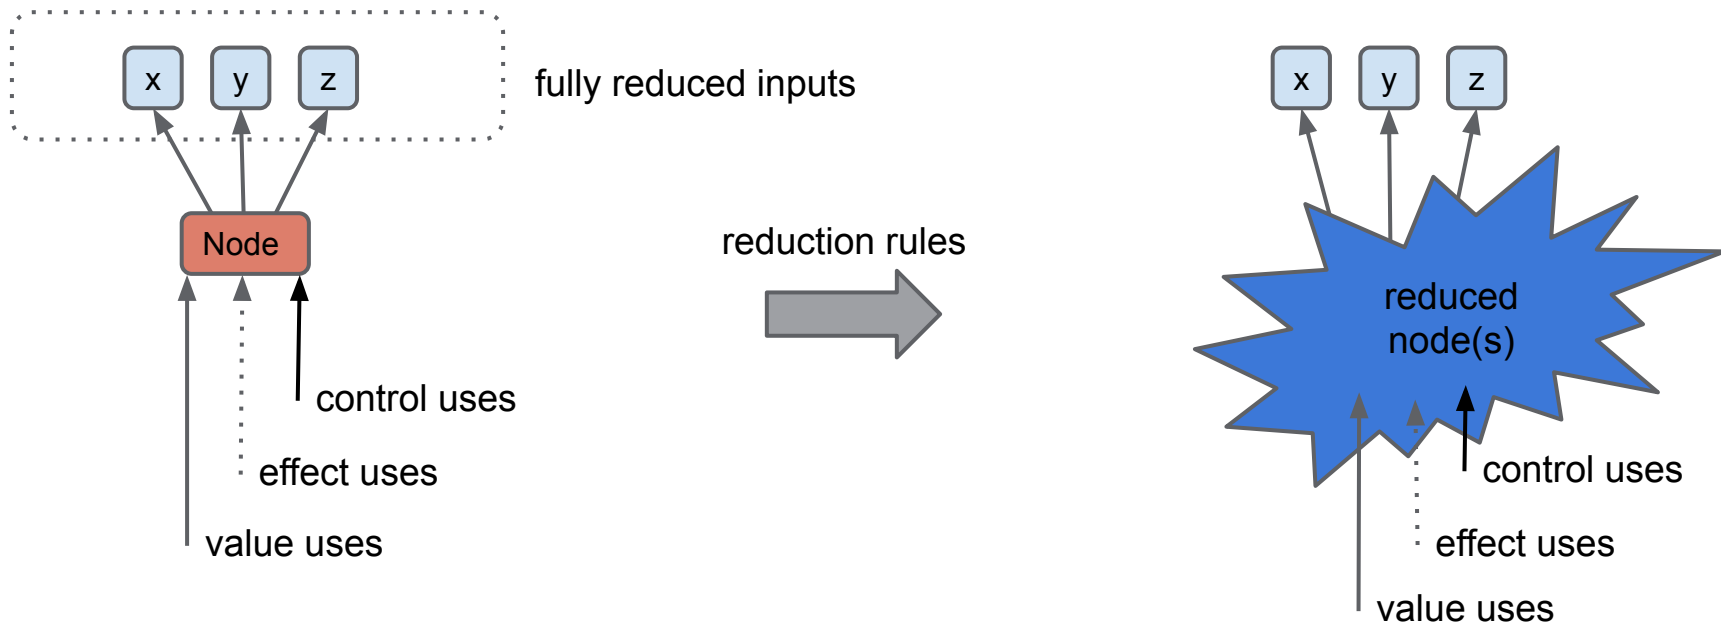






# Control Optimization as Reduction



# Reduction as Top-down Graph Rewriting



# Iterative Reduction (recursion with explicit stack)





Reduce top of stack when all its inputs are reduced.

Pop the stack after applying reduction rules.

Applies reduction to each node once in the optimal order.

# Iterative Reduction (in the presence of cycles)





node stack

Reduce top of stack when all its inputs are either reduced ***or are themselves on the stack.***

When a node is successfully reduced, revisit any uses that were partially reduced due to cycles.

Computes a fixpoint over reduction rules.

# Iterative Reduction (in the presence of cycles)





Reduce top of stack when all its inputs are either reduced ***or are themselves on the stack.***

When a node is successfully reduced, revisit any uses that were partially reduced due to cycles.

Computes a fixpoint over reduction rules.

# Iterative Reduction (in the presence of cycles)





Reduce top of stack when all its inputs are either reduced or are themselves on the stack.

When a node is successfully reduced, ***revisit any uses that were partially reduced due to cycles.***

Computes a fixpoint over reduction rules.

# Iterative Reduction (in the presence of cycles)



node stack


Reduce top of stack when all its inputs are either reduced or are themselves on the stack.

When a node is successfully reduced, revisit any uses that were partially reduced due to cycles.


**Computes a fixpoint over reduction rules.**

# Lowering to Machine


- JavaScript:



- Intermediate:




- Machine:







# Lowering to Machine




expand and optimize  
JS\* and Simplified\*  
nodes




# Floating Control




lower and  
expand fast  
case



# Floating Control






lower and  
expand fast  
case




# Scheduling the Sea of Nodes

- Sea of nodes expresses many possible legal orderings of code
  - Many possible CFGs
  - Many possible assignment of nodes to CFG blocks
  - Many possible orderings within basic blocks
- What is the most efficient order and placement?
  - Depends on control dominance, loop nesting, register pressure
- Outcome: traditional CFG
  - Traditional code generation and register allocation can take over



# Scheduling the Sea of Nodes (sketch)



# Scheduling the Sea of Nodes (sketch)






Place fixed  
nodes  
(phis, params)






Dominator  
Tree

# Scheduling the Sea of Nodes (sketch)






# Scheduling the Sea of Nodes (sketch)











# Scheduling the Sea of Nodes (sketch)






# Scheduling the Sea of Nodes (sketch)





# Scheduling the Sea of Nodes (sketch)




# Scheduling the Sea of Nodes (sketch)





# Scheduling the Sea of Nodes (sketch)




Schedulable nodes




# Scheduling the Sea of Nodes (sketch)



# Scheduling the Sea of Nodes (sketch)




Schedulable  
nodes




Dominator  
Tree

# Scheduling the Sea of Nodes (sketch)




A fully scheduled graph is exactly the same as a CFG.





# Scheduling the Sea of Nodes (sketch)



Control Flow Graph


Placement is important!

Hoist code out of loops if possible

Place code as late as possible


Minimize register pressure

Eliminate redundancy



Dominator Tree


# Where is Loop Invariant Code Motion?



Control Flow Graph


Nowhere!

Scheduling subsumes loop  
invariant code motion.




Dominator  
Tree

# Instruction Selection (theory)

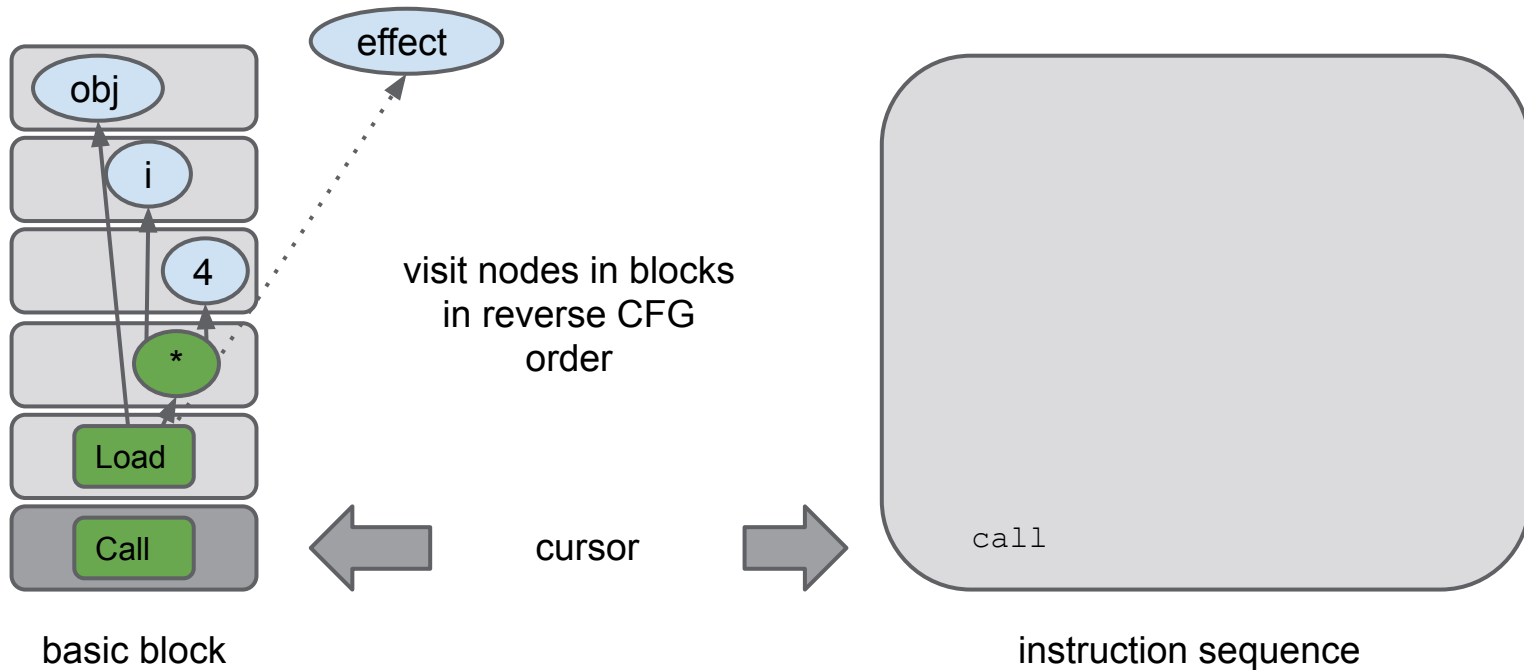


maximal  
munch

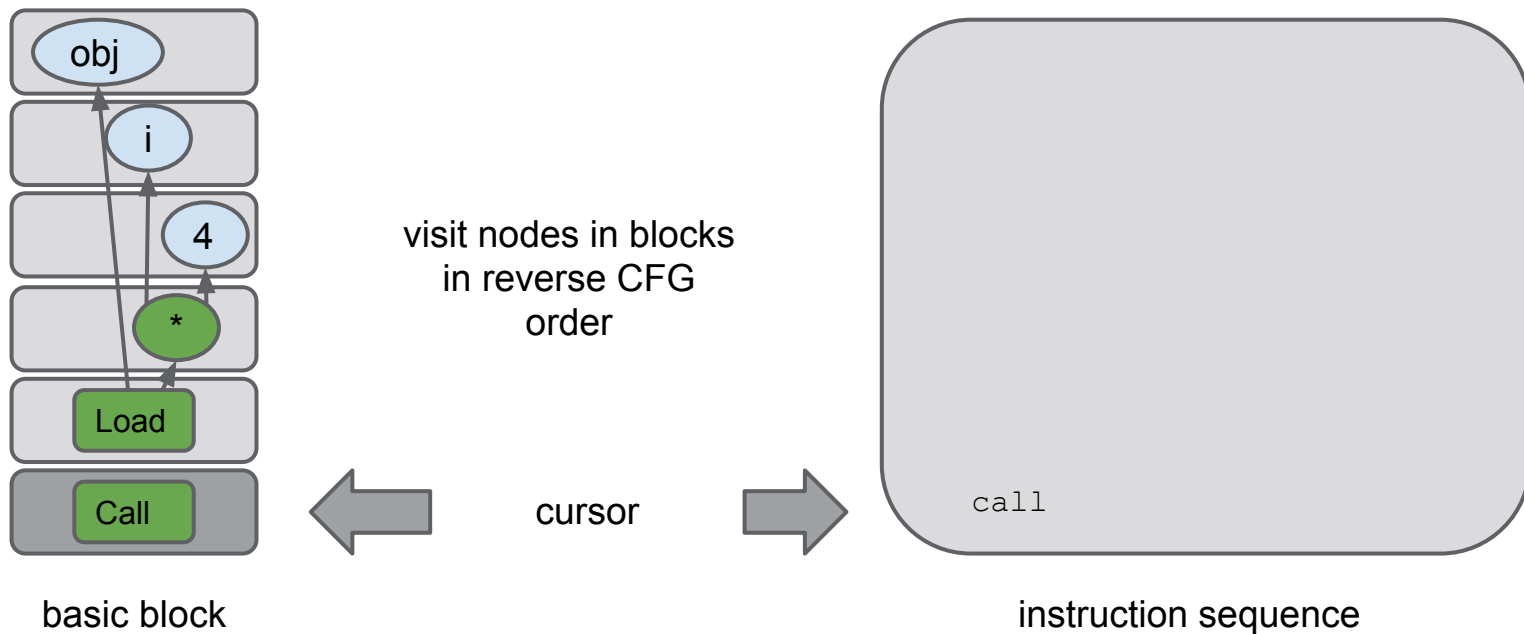


```
mov %r0, [%r1 + %r2 * 4]
```

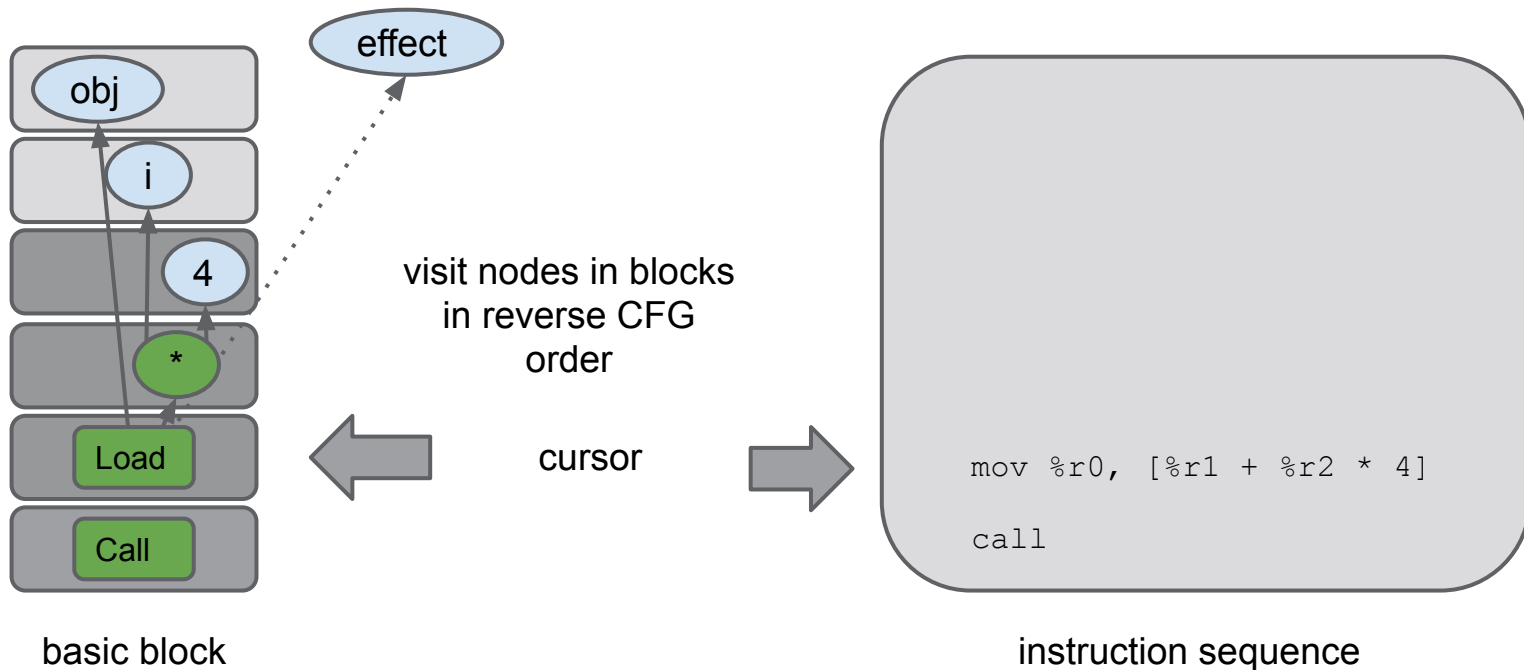
# Instruction Selection (practice)



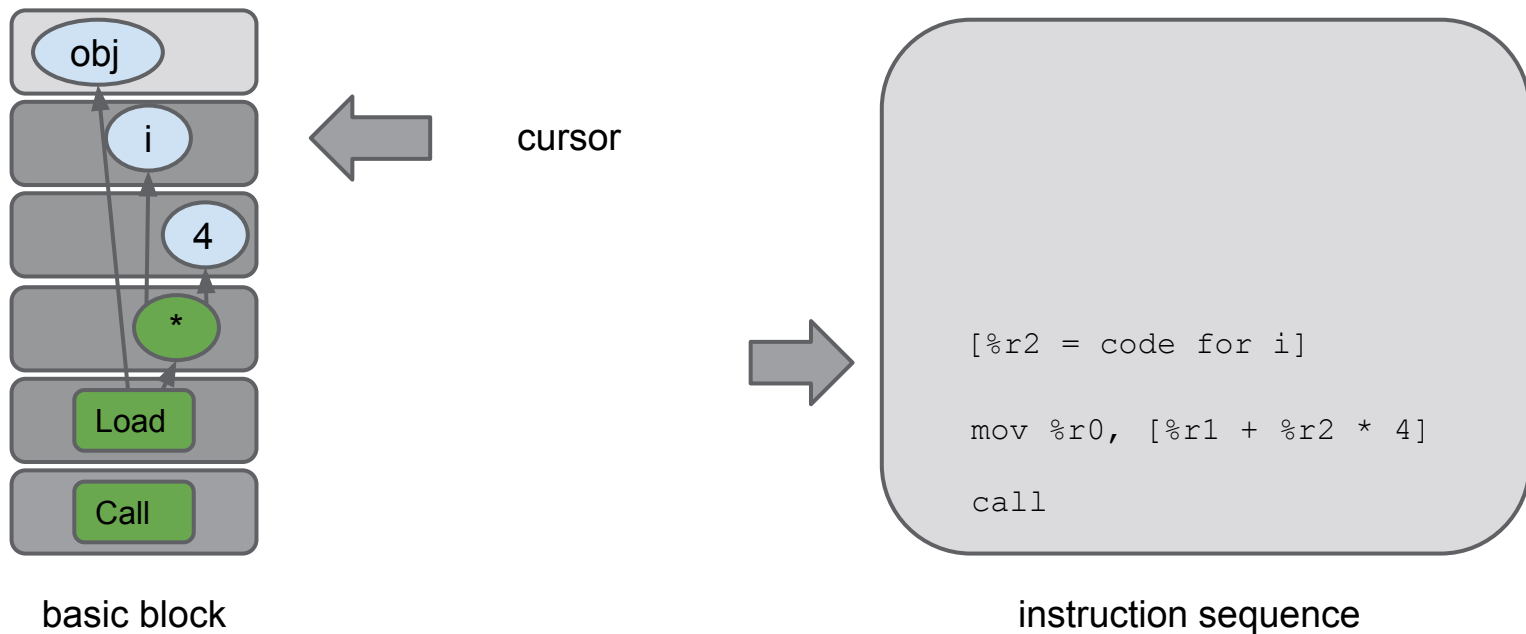
# Instruction Selection (practice)



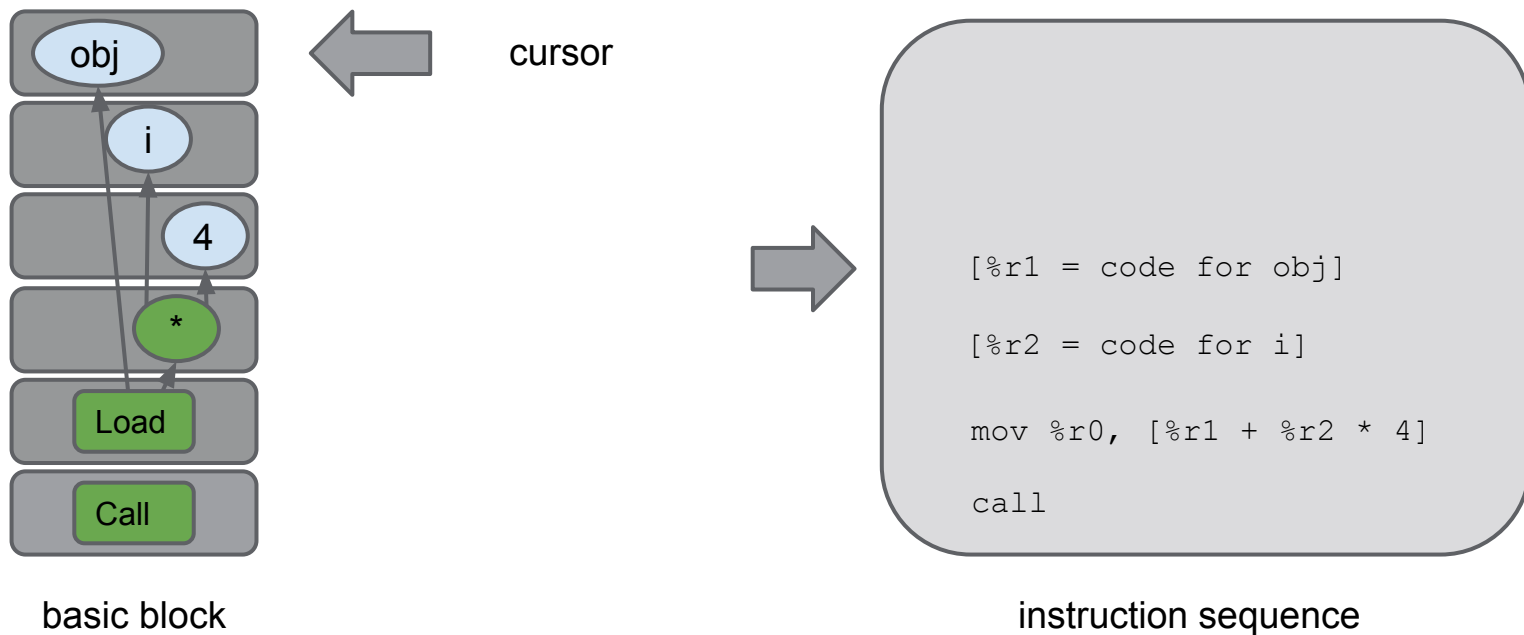
# Instruction Selection (practice)



# Instruction Selection (practice)



## Instruction Selection (practice)





# Register Allocation

TurboFan uses a linear scan allocator with live range splitting to assign registers and insert spill code.

SSA form can be deconstructed before or after register allocation with explicit moves.

```
[%r1 = code for obj]
```

```
[%r2 = code for i]
```

```
mov %r0, [%r1 + %r2 * 4]
```

```
call
```

instruction sequence

# Register Allocation

TurboFan uses a linear scan allocator with live range splitting to assign registers and insert spill code.

SSA form can be deconstructed before or after register allocation with explicit moves.

The result of register allocation is to replace uses of virtual registers with real registers and to insert spill code between instructions.

```
mov [sp + 12], %eax
```

```
[%eax = code for obj]
```

```
[%ebx = code for i]
```

```
mov %ecx, [%eax + %ebx * 4]
```

```
call
```

instruction sequence


# Testability

- Unit testing
  - Basic data structures, traversal algorithms, lowering, graph building, graph transformations, type relations, instruction selection, spill code insertion, register allocation, code generation, assemblers
- Integration testing
  - Run multiple optimization passes together, create specific graphs explicitly, try all combinations of arithmetic + generate and run code
- Performance tracking
  - Microbenchmarks, benchmark suites
- Fuzzing
  - Randomly mutate JavaScript source and feed it compiler

# Status

- TurboFan is now in beta testing in Chrome 41
  - Initially enabled for asm.js
- 6 Fully supported platforms
  - ia32, x86-64, arm, arm64, mips, mips64
  - 2000-3000 lines per platform (vs 13000-16000 for CrankShaft)
- Improves Octane zlib benchmark 30-45%
- Most Emscripten benchmarks 10-25% faster
- Working on “general” JavaScript
  - Goal: all of ES6 within a couple of months

# Thank You!



Ben L. Titzer

Michael Starzinger

Daniel Clifford

Benedikt Meurer

Dan Carney

Andreas Rossberg

Jaroslav Sevcik

Sven Panne

Sigurd Schneider

Georg Neis